

## Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader\_weights(), grader\_sigmoid(), grader\_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn import linear_model

Creating custom dataset

In [2]: # Please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                           n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_classification.html) for more details

In [3]: X.shape, y.shape
Out[3]: ((50000, 15), (50000,))

In [4]: len(X)
Out[4]: 50000

Splitting data into train and test

In [5]: #Please don't change random state
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)

In [6]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
Out[6]: ((37500, 15), (37500,)), (12500, 15), (12500,))

In [7]: X_train[0]
Out[7]: array([-0.57349184, -0.19015688, -0.06584143, -0.86909562, -2.80927706,
               -1.43345052,  0.35862361,  0.24627836, -2.25803168, -0.87761289,
                2.31023199, -0.3484947 , -2.2575668 , -1.93628665,  1.65242231])
```

## SGD classifier

```
In [8]: # alpha : float
# Constant that multiplies the regularization term.
# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.
clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15,
                                penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html)

Out[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                     early_stopping=False, epsilon=0.1, eta0=0.0001,
                     fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                     loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                     penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                     tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)

In [9]: clf.fit(X=X_train, y=y_train) # fitting our model

-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.45552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.05 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.06 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.07 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819025, T: 300000, Avg. loss: 0.378856
Total training time: 0.08 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837005, T: 337500, Avg. loss: 0.378585
Total training time: 0.09 seconds.
-- Epoch 10
Norm: 1.09, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.10 seconds.
Convergence after 10 epochs took 0.10 seconds

Out[9]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                     early_stopping=False, epsilon=0.1, eta0=0.0001,
                     fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                     loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                     penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                     tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)

In [10]: clf.coef_, clf.coef_.shape, clf.intercept_
clf.coef_ will return the weights
clf.coef_.shape will return the shape of weights
clf.intercept_ will return the intercept term

Out[10]: (array([[ -0.42366092,  0.18547565, -0.14859036,  0.34144407, -0.2091067 ,
                0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.10084126,
                0.1705191 ,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
         (15,))
array([[-0.8531383]])
```

## Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

- We will be giving you some functions, please write code in that functions only.
- After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight\_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)

log loss =  $-1 * \text{frac}(1/(n)) \text{Sigma\_for each } Y_i, Y_{\text{pred}}(Y \log(1/Y_{\text{pred}}) + (1 - Y) \log(1 - Y_{\text{pred}}))$

- for each epoch:
  - for each batch of data points in train: (keep batch size=1)
    - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)

$\text{Sdw}(\text{w}) = x_n y_n - o((w^{(0)})(T) x_n + b^{(0)}) - \text{frac}(N)(N w^{(0)})$

- Calculate the gradient of the intercept (write your code in `def gradient_db()`) [check this](#)

$\text{Sdb}(\text{w}) = y_n - o((w^{(0)})(T) x_n + b^{(0)})$

- Update weights and intercept (check the equation number 32 in the above mentioned pdf:  $\text{Sdw}(\text{w}^{(t+1)}) - \text{w}^{(t)} + o(\text{dw}^{(t)})$ )
- $\text{Sb}(\text{w}^{(t+1)}) - b^{(t)} + o(\text{db}^{(t)})$
- calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
- And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
- append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [11]: def initialize_weights(dim):
''' In this function, we will initialize our weights and bias'''
# initialize the weights to zeros array of (dim,1) dimensions
# you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
# initialize bias to zero
w=np.zeros_like(X_train[0])
b=0
return w,b

In [12]: dim=X_train[0]
w,b = initialize_weights(dim)
print('w =',w)
print('b =',str(b))

w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function -1

```
In [13]: dim=X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
assert((len(w)==len(dim)) and b==0 and np.sum(w)==0)
return True
grader_weights(w,b)

Out[13]: True
```

Compute sigmoid

```
$sigmoid(z) = 1/(1+exp(-z))$

In [14]: def sigmoid(z):
''' In this function, we will return sigmoid of z'''
# compute sigmoid(z) and return
sig= 1/(1 + np.exp(-z))
return sig
```

Grader function -2

```
In [15]: def grader_sigmoid(z):
val=sigmoid(z)
assert(abs(val-0.8079707797778223)
return True
grader_sigmoid(z)

Out[15]: True
```

Compute loss

log loss =  $-1 * \text{frac}(1/(n)) \text{Sigma\_for each } Y_i, Y_{\text{pred}}(Y \log(1/Y_{\text{pred}}) + (1 - Y) \log(1 - Y_{\text{pred}}))$

```
In [16]: import math
from math import log10

def logloss(y_true,y_pred):
'''In this function, we will compute log loss'''
n=len(y_true)
loss=0
for i in range(len(y_true)):
loss += (-1)*(1/n)*np.sum((y_true[i]*(math.log10( y_pred[i])))+(1-y_true[i]*(math.log10(1-y_pred[i]))))

return loss
```

Grader function -3

```
In [17]: def grader_logloss(true,pred):
loss=logloss(true,pred)
assert(loss==0.8764490402910309)
return True
true=[1,1,0,1,0]
pred=[0.9,0.0,0.0,0.1,0.0,0.2]
grader_logloss(true,pred)

Out[17]: True
```

Compute gradient w.r.to w'

$\text{Sdw}(\text{w}) = x_n y_n - o((w^{(0)})(T) x_n + b^{(0)}) - \text{frac}(N)(N w^{(0)})$

```
In [18]: def gradient_dw(x,y,w,b,alpha,N):
'''In this function, we will compute the gradient w.r.to w'''
z=np.dot(w,T,x) + b
dw= x*(y-sigmoid(z) - (alpha/N) * w)
return dw

In [19]: y_train
Out[19]: array([0, 0, 0, ..., 1, 0, 0])
```

Grader function -4

```
In [20]: def grader_dw(x,y,w,b,alpha,N):
grad_dw=gradient_dw(x,y,w,b,alpha,N)
assert(np.sum(grad_dw)==2.613689585)
return True
grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
                -2.81434437, -0.86771071, -0.04073287, 0.84827878, 1.99451725,
                3.07152472, 0.01451075, 2.01002086, 0.07373904, -5.54500692])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)

Out[20]: True
```

Compute gradient w.r.to b'

$\text{Sdb}(\text{w}) = y_n - o((w^{(0)})(T) x_n + b^{(0)})$

```
In [21]: def gradient_db(x,y,w,b):
'''In this function, we will compute gradient w.r.to b'''
db = y - sigmoid( ( np.matmul(w,x) ) + b )
return db
```

Grader function -5

```
In [22]: def grader_db(x,y,w,b):
grad_db=gradient_db(x,y,w,b)
assert(grad_db==0.5)
return True

grad_x=np.array([-2.07864835, 3.31604252, -0.79104357, -3.87045546, -1.14783286,
                -2.81434437, -0.86771071, -0.04073287, 0.84827878, 1.99451725,
                3.07152472, 0.01451075, 2.01002086, 0.07373904, -5.54500692])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(X_train)
grader_db(grad_x,grad_y,grad_w,grad_b)

Out[22]: True
```

Implementing logistic regression

$w^{(t+1)} = w^{(t)} + o(dw^{(t)})$

$b^{(t+1)} = b^{(t)} + o(db^{(t)})$

```
In [23]: def train(X_train,y_train,X_test,y_test,epochs,alpha,eta0):
''' In this function, we will implement logistic regression'''
# eta0 is learning rate
# implement the code as follows
# initialize the weights (call the initialize_weights(X_train[0]) function)
# for every epoch
# for every data point(X_train,y_train)
# compute gradient w.r.to w (call the gradient_dw() function)
# compute gradient w.r.to b (call the gradient_db() function)
# update w,b
# predict the output of x_train(for all data points in X_train) using w,b
# compute the loss between predicted and actual values (call the loss function)
# store all the train loss values in a list
# predict the output of x_test(for all data points in X_test) using w,b
# compute the loss between predicted and actual values (call the loss function)
# store all the train loss values in a list
# you can also compare previous loss and current loss, if loss is not updating then
stop the process and return w,b
w,b=initialize_weights(X_train)
N=len(X_train)
loss_train=[]
loss_test=[]
for epoch in range(epochs):
for i in range(len(X_train)):
dw=gradient_dw(X_train[i],y_train[i],w,b,alpha,N)
db=gradient_db(X_train[i],y_train[i],w,b)
w = w + eta0 * dw
b = b + eta0 * db
z_train=np.dot(X_train,w)+b
y_pred_train= sigmoid(z_train)
y_pred_test= sigmoid(z_test)

loss_train=logloss(y_train,y_pred_train)
loss_test=logloss(y_test,y_pred_test)
loss_train.append(loss_train)
loss_test.append(loss_test)
return w,b,loss_train,loss_test
```

```
In [24]: alpha=0.0001
eta0=0.0001
N=len(X_train)
epochs=75
w,b,train_loss,test_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0)
```

```
In [25]: w
Out[25]: array([[ -0.42979252,  0.19303522, -0.14846993,  0.33809367, -0.22128249,
                0.569949 , -0.44051863, -0.08909394,  0.22182953,  0.17382971,
                0.19874852, -0.0005843 , -0.08133412,  0.33909013,  0.02298796]])
```

```
In [26]: b
Out[26]: -0.8922527269940091
```

```
In [27]: train_loss
Out[27]: [0.754574843371859,
0.16867157942532082,
0.16639167986531958,
0.16536827532926193,
0.16485707456106383,
0.1645882001023163,
0.16444271321223278,
0.16436263614118496,
0.1643180694539515,
0.16429307373049495,
0.16427097430077493,
0.16427098545161656,
0.16426644190475895,
0.16426304191012098,
0.16426236407945032,
0.1642615119026574,
0.1642610212937372,
0.16426073527358003,
0.1642605693872455,
0.16426047215030304,
0.16426041409604803,
0.1642603804167069,
0.1642603597246379,
0.1642603470619127,
0.1642603391900077,
0.16426033421015734,
0.16426033099990248,
0.16426032888966322,
0.1642603274752386,
0.16426032650929026,
0.1642603258381035,
0.1642603256445398,
0.1642603250250840,
0.1642603247900234,
0.16426032460168455,
0.1642603244710028,
0.16426032437275702,
0.16426032430009743,
0.1642603242509896,
0.1642603242054804,
0.1642603241750857,
0.16426032415224126,
0.16426032413505123,
0.16426032412209499,
0.1642603241123221,
0.16426032410496603,
0.16426032409940797,
0.1642603240952125,
0.16426032409204309,
0.16426032408964863,
0.16426032408784247,
0.16426032408647456,
0.16426032408544307,
0.16426032408466645,
0.164260324084007643,
0.1642603240836309,
0.164260324083295,
0.1642603240829653,
0.1642603240826506,
0.16426032408270355,
0.1642603240825944,
0.1642603240825112,
0.16426032408244906,
0.16426032408240224,
0.1642603240823657,
0.16426032408233895,
0.16426032408231786,
0.1642603240823004,
0.16426032408229024,
0.1642603240822825,
0.1642603240822825,
0.16426032408227592,
0.16426032408227031,
0.16426032408226723,
0.164260324082264,
0.16426032408226254]
```

```
In [28]: test_loss
Out[28]: [0.1759547441481585,
0.16939931352758589,
0.16720591191229046,
0.16621717797172324,
0.1657195946374875,
0.1654595719820607,
0.1653113002208009,
0.16523116855206854,
0.16518095100901304,
0.1651604056163819,
0.16514582031758648,
0.16513739038632227,
0.16512951007837477,
0.16512967662910438,
0.16512800522364235,
0.1651270142478074,
0.16512642053988932,
0.1651260060779014,
0.16512583901790344,
0.16512570062122825,
0.16512561260498806,
0.16512555557494782,
0.16512551790645463,
0.16512549254869247,
0.16512547156554420,
0.16512546304812392,
0.16512544544745603,
0.16512544032740512,
0.16512543430708030,
0.1651254317737316,
0.1651254311712407,
0.1651254311666023,
0.1651254311630947,
0.16512543116044787,
0.16512543115844633,
0.1651254311569348,
0.16512543115579245,
0.16512543115492825,
0.1651254311541494,
0.16512543115378228,
0.16512543115340966,
0.16512543115312822,
0.16512543115291567,
0.16512543115275427,
0.1651254311526331,
0.16512543115254108,
0.16512543115247136,
0.1651254311524185,
0.16512543115237385,
0.16512543115234904,
0.16512543115232567,
0.16512543115230097,
0.16512543115229633,
0.16512543115228726]
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of  $10^{-3}$

```
In [29]: # these are the results we got after we implemented sgd and found the optimal weights and in
# train
w=clf.coef_, b=clf.intercept_

Out[29]: (array([[ -0.00642561,  0.00755957,  0.00012042, -0.0033604 , -0.01309579,
                0.00976321,  0.00724319,  0.00418419,  0.01256633, -0.00781155,
                0.00169611, -0.00408345, -0.00173043,  0.000856212,  0.000320751]]),
         array([-0.03911443]))
```

Plot epoch number vs train\_loss

- epoch number on X-axis
- loss on Y-axis

```
In [30]: def pred(w, X):
N = len(X)
predict = []
for i in range(N):
z = np.dot(w,[1] + b
res=sigmoid(z)
if res.any() >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
else:
predict.append(0)
return np.array(predict)

print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))

1.6978933333333335
1.6986000000000001
```

```
In [32]: import matplotlib.pyplot as plt
plt.plot(range(epochs), train_loss, label='train_loss')
plt.plot(range(epochs), test_loss, label = 'test_loss')
plt.legend()
plt.xlabel("epoch")
plt.ylabel("Loss")
plt.title("train and test loss")

plt.show()
```

